

Exemplars Move Retrieval, Instructions Move Rejection: An Anatomy of Inference-Time Interventions for Closed-Book Knowledge Base Construction

Takashi Sasaki

Acroquest Technology Co., Ltd.
sasaki@acroquest.co.jp

Abstract

We describe our system for the AKBC 2026 shared task on closed-book knowledge base construction. Our final system is a single pretrained-only backbone, Gemma-3-27B-pt, queried entirely in a raw text-completion format, with 20 candidates per query, relation-specific gated aggregation, and no arbiter; it uses 27B of the 32B parameter budget and scores a row-averaged Macro F1 of 0.650 on validation and 0.6493 on the official test set, against 0.313 for the official baseline. Its predecessor—two deliberately different prompts over Mistral Small 3.2 24B with a cross-family 8B arbiter, 31.8B in total—scores 0.6247. Most of the paper presents a mechanism study over three backbones and roughly thirty controlled experiments, organized around three results. First, how reliably a backbone recalls the value stored for an entity label dominated every intervention we tried: each backbone change moved the score more than all our prompt, sampling, and selection experiments combined, and the winning backbone is a pretrained checkpoint in a completion format, not an instruct model in a chat format. Second, exemplar style moves retrieval and instruction text does not—two in-prompt exemplars demonstrating a verbalized recall move the answer mode by +6 rows, while the same request as an instruction moves nothing—but instructions do move *rejection*, the discarding of already-generated candidates. Third, validation gains transferred when the change was gated (applied only where a condition holds) *and* the gate consulted information from outside the model; a mechanism-derived but ungated rescoring gained +2 validation rows and lost a test row.

1 Introduction

The LM-KBC shared task series (Singhanian et al., 2023; Kalo et al., 2024) probes how much factual knowledge language models can turn into knowledge base entries: given a subject and a relation,

a system must output *all* correct object values—possibly none, one, or hundreds—under closed-book conditions: no retrieval, no external corpora, no fine-tuning, and at most 32B neural parameters at inference.

Under these conditions the only controls are at inference time: model choice, prompts, sampling, and aggregation over samples. Our development proceeded in three stages: a Qwen3.5-27B thinking system (test 0.571), a Mistral Small 3.2 system with a second prompt and a cross-family 8B arbiter (0.6247), and a single pretrained-only Gemma-3-27B backbone in a completion format (0.6493). Alongside the system we report the controlled experiments behind each decision, including submissions whose test scores give us what validation cannot: isolated ablations on the hidden set.

1. **Backbone recall reliability dominates.** A 12-row probe (30 GPU-minutes per model) ranked backbones consistently with downstream quality in every deployment decision we evaluated, and both backbone changes beat all our prompt and selection work combined (Sec. 4.1).
2. **Answers commit early** to the first concrete value they generate (Sec. 4.2), across backbones.
3. **Exemplar style moves retrieval; instruction text does not** (Sec. 4.3); instructions move rejection. The effect is conditional on recall reliability and on post-training (Sec. 4.4).
4. **Within one model we found no usable selection signal;** across prompts and families we did (Sec. 4.5).
5. **Validation gains transfer when the change is gated and the gate looks outside the model** (Sec. 4.6).

6. **The residual errors are consistent with incorrect values learned from training data**, and are associated with the instruct-plus-chat setting (Sec. 4.7).

2 Task and Constraints

The 2026 task covers six relations over 478 validation and 477 test subject–relation pairs: two set-valued relations that may be empty (`countryLandBordersCountry`, `companyTradesAtStockExchange`), one single-valued string relation that may be empty (`personHasCityOfDeath`), one large-set relation (`awardWonBy`), and two numeric relations (`hasArea`, `hasCapacity`) scored at 5% relative tolerance. Gold values come from Wikidata (areas P2046, capacities P1083), with alias lists accepted for strings. We report the official “All Relations” score throughout: a Macro F1 averaged over rows, not relations, so the four 100-row relations dominate it. Systems are closed-book and capped at 32B inference parameters.

3 System Description

3.1 Final system

The final system generates with Gemma-3-27B-pt (the pretrained checkpoint, Q8_0, served by `llama.cpp` at a pinned revision), one model and no arbiter, 27B of the 32B budget. Every relation is queried in a *raw text-completion format*—a record- or infobox-style prefix that the model continues, with no chat template and no instruction turn—with 20 candidates per query at temperature 0.6, top- p 0.95 and deterministic per-candidate seeds, as in the earlier systems.

Aggregation is relation-specific and, in every case, *gated*: a rule applies only where an explicit support condition holds and falls back to a default otherwise. Numeric candidates are clustered at 5% relative tolerance and `dominant_cluster` predicts the median of the best-supported cluster. `personHasCityOfDeath` takes a surface majority with a minimum-vote gate, predicting empty below it. `companyTradesAtStockExchange` folds surface variants—case variants and acronyms that expand to a co-occurring long form (NYSE → New York Stock Exchange)—into a single count, repairing vote splits that would drop majority answers below its frequency threshold (≥ 0.5), and abstains on an empty majority. `countryLandBordersCountry` and `awardWonBy`

use recall-oriented thresholds (≥ 0.3 , ≥ 0.1) with the same abstention.

Reproducibility. Code, configurations, builders, the step-by-step recipe (`docs/V4_REPRODUCE.md`), and per-experiment records are at <https://github.com/acrosasaki-rgb/llm-knowledge-task-2026>.

Model revisions, container digests, and the dataset commit are pinned; each prediction file is a deterministic function of persisted candidate sets and a builder script (SHA-256 verified).

3.2 Predecessor system

The mechanism experiments of Sec. 4, and all our per-component test measurements, come from the Mistral-era system that preceded it. It samples 20 candidates per prompt from Mistral Small 3.2 24B under *two* prompts of deliberately different style for `hasArea`: a **recall prompt** with no few-shot examples but two in-prompt exemplars demonstrating a labeled recall format (Entity:, Memory: “I recall that the Wikipedia infobox lists X’s total area as Y”, Unit:, Answer:), and a **grounding prompt** with five bare-JSON few-shot examples and an instruction naming the recall target (Wikidata P2046/P1083). `hasCapacity` adds a third prompt at zero additional parameter-count cost: the same 24B model queried as raw text completion. When two prompts disagree by more than 5%, one aggregated arbiter decision—the mode over 20 samples of an unrelated-family Llama-3.1-8B—breaks the tie, switching to the secondary answer only if the arbiter matches it and not the primary one. The arbiter never introduces a value neither prompt proposed, but an incorrect switch can still lose a row. Total: $24\text{B} + 8\text{B} \approx 31.8\text{B}$.

3.3 Results

Three contributions are separable on test: the Qwen→Mistral backbone change (+0.024, rest of the pipeline unchanged), the Mistral-era pipeline interventions (+0.015, v3→v5), and the Mistral→Gemma-pt backbone change (+0.025, v11→v12). A further +0.015 came from following the dataset re-keying, which corrects what is scored rather than improving the system. Per-relation test scores for v12: `borders` 0.954, `company` 0.842, `area` 0.810, `city` 0.490, `award` 0.336, `capacity` 0.276. The new backbone gains most over v11 on `company` ($0.778 \rightarrow 0.842$) and `award` ($0.279 \rightarrow 0.336$) and loses on `city` ($0.510 \rightarrow 0.490$); the two agree on 32% of `hasArea` and 29% of `hasCapacity` rows.

System	cap	val	test
Baseline (Qwen3.5-9B)	✓	0.313	—
Qwen3.5-27B thinking	✓	0.588	0.571
Mistral 3.2 grounding	✓	0.595	0.595
v3: + recall prompt, area			
arbitration, merge	✓	0.629	0.604
v4: + cap. arbitration	✓	0.633	0.606
v5: + acronym merging	✓	0.635	0.610
v6: + second arbiter	×	0.637	0.610
v7: + city merge	×	0.660	0.608
v8: v6 + re-keying*	×	0.637	0.625
v10: + bridge, cap fix	✓	0.641	0.623
v11: — bridge	✓	0.637	0.625
v12: Gemma-3-27B-pt	✓	0.650	0.649

Table 1: System progression (row-averaged Macro F1, three decimals). Each row is its predecessor plus the named change, so the test column doubles as a sequence of isolated ablations on the hidden set (Sec. 4.6); *cap* marks configurations within the 32B limit. *Following the August subject-disambiguation release, which renamed 15 test subjects and dropped two rows (477→475); v3–v7 score the renamed rows as zero. v9 (the rescoring on the non-compliant configuration) is omitted for space.

Model (12-row probe)	any	mode
Mistral Small 3.2 Instruct	11	8
Mistral Small 3.1 Base	10	7
Mistral Small 3.1 Instruct	9	4
Llama-3.1-8B Instruct	7	4
Aya Expanse 32B	3	1
Qwen3-8B	1	0

Table 2: The recall probe: 12 hasArea rows stratifying the Qwen-era system’s failures, 20 candidates each; *any* = rows with a correct candidate, *mode* = rows whose modal cluster is correct. Qwen3.5/3.6-27B, Gemma 4 31B, GLM-Z1-32B and llm-jp-4-32b (not shown) sit between Llama-8B and Aya. It costs ~30 GPU-minutes per model.

4 Anatomy of Inference-Time Interventions

All experiments change one stated condition; numeric testbeds are hasArea and hasCapacity (100 validation rows × 20 candidates per prompt). **Oracle** is the number of rows whose candidate set contains at least one within-tolerance value; the **mode** row count is what dominant_cluster selects. The **anchor** is the first concrete value for the queried quantity in a candidate.

4.1 Backbone Recall Reliability Dominates

Our Qwen-era system plateaued after nine intervention experiments (prompts, temperature, out-

put routes, thinking budgets, judge models, log-probability confidence) whose combined deployable gain was under +0.01. A 12-row probe of its wrong rows across candidate models (Table 2) showed Mistral Small 3.2 answering 11 of the 12 somewhere in its candidate set; regenerating the pipeline on Mistral with nothing else changed moved test 0.571→0.595.

The second change was larger and reversed an earlier decision. Gemma 4 31B Instruct, queried through its chat template, had the weakest hasArea oracle we measured (43/100) and we closed that lineage early. Probing the *pretrained* Gemma-3-27B checkpoint in a completion format instead gave the best profile of any configuration we tested—oracle 94, mode 84 on hasArea, borders 0.986—and a system built on it took test 0.625 → 0.649. It also supplies values no prompt reached on Mistral: secondary stock listings that five Mistral prompt variants returned with zero votes (ČEZ on Warsaw, Questerre on Oslo, Aisin on Nagoya) appear with 7–13 votes out of 20. This was the dominant empirical factor in our experiments: closed-book quality is set by which label-to-value pairings pretraining left intact and by whether the prompting format can reach them.

4.2 Answers Commit Early

In persisted reasoning traces of the Qwen-era system, the answer value first appears at a median 17% of the trace, in 95–98% of candidates inside a retrieval template (“Commonly cited: X”); the remaining 80% of the trace never dislodges it. The verification text that follows is consistent with post-hoc justification rather than independent checking: fabricated Wikidata Q-IDs and cross-language “confirmations” agree with whatever the anchor was, and wrong candidates carry *more* of this apparatus than correct ones (6.9 vs. 4.2 source mentions). Early commitment replicates on Gemma 4 31B and Qwen3.6; the confabulated verification does not. Extended thinking is thus an *exploration* aid: it raises the oracle where recall is unreliable (Qwen +9 rows, Gemma +7) and has no measurable effect on Mistral or Magistral.

4.3 Exemplar Style Moves Retrieval; Instruction Text Does Not

The single largest prompt effect we found—hasArea mode 72 → 78, transferring to test as 0.70 → 0.75—looked at first like an instruction effect: a step-structured expert prompt, adapted

hasArea prompt (Mistral 3.2)	Few-shot examples	
	bare-JSON $\times 5$	none
grounding instruction	71 (o. 89)	41 (o. 59)
recall-format text	72 (o. 83)	78 (o. 84–89)

Table 3: The 2×2 comparison: mode-correct rows out of 100 (oracle in parentheses). The recall-format text carries two exemplars inside the instruction, so the bottom-right cell keeps two exemplars while the top-right has *none* and collapses to 41. At fixed example style the instruction moves the result by ± 1 ; the example style moves it by $+6 / -30$.

from a style published by another participant, beat our grounding instruction by 6 rows. The 2×2 comparison (Table 3) shows the instruction text is nearly inert. What matters is the *exemplars*: with two in-prompt exemplars demonstrating a one-sentence verbalized recall (“I recall that the Wikipedia infobox lists Luxembourg’s total area as 2586.4 km²”), 99% of candidates reproduce the verbalization and the mode moves $+6$; with five bare-JSON few-shot exemplars, 0% verbalize—the same procedure requested as instruction text is silently ignored—and the mode does not move. Iterating the wording (persona, fixed phrase, provenance vocabulary) changed nothing, while reproducing the labeled structure with our own labels recovered the full effect.

Instructions move rejection. Instruction text was not inert everywhere: it moved the *discarding* of already-generated candidates. Verifying each awardWonBy candidate with three yes/no votes and dropping only unanimous-no candidates raised precision $0.241 \rightarrow 0.280$ and F1 $0.172 \rightarrow 0.186$; the yes votes carried no information, so the usable signal was one-sided. A digit-check step inside an otherwise inert step-by-step procedure likewise cut votes for order-of-magnitude twin values from 38 to 32 with the oracle unchanged. Instructions act after retrieval, on which candidates survive.

4.4 The Exemplar Effect Is Conditional

Verbalized recall is a *focusing* intervention—it sharpens the anchor distribution around the most frequently recalled value (oracle drops $89 \rightarrow 84$ as the mode improves). Its sign therefore follows how reliable the model’s recall already is, in mirror image to thinking: thinking explores, gaining $+9$ rows where recall is unreliable (Qwen) and nothing where it is reliable (Mistral); verbalized recall does the opposite, zero or negative on Qwen and

Component	val	test	transfers
recall exemplars (§4.3)	+6	+5	✓
company folding (§3)	+1	+2	✓
8B arbitration (§4.5)	+2	+1	✓
capacity prompt (best of 6)	+2	−2	×
second arbiter (increment)	+1	0	×
tolerance rescoring	+2	−1	×
city two-stage merge	+11	−1	×

Table 4: Row-level validation and test deltas for deployed components, from the isolated ablations of Table 1. The split follows neither effect size (our largest validation gain transferred worst) nor the presence of a mechanism-level rationale (the rescoring has one).

on Mistral’s weakest relation (hasCapacity oracle $60 \rightarrow 47$) but $+6$ on hasArea. In Magistral (Mistral’s RL-tuned sibling) the effect disappears: it copies the verbalized form at 99.9% compliance—in its thinking channel, in a scripted thinking channel, and in the answer body—yet its anchor distribution does not move in any of the three (70–72), suggesting that RL post-training may decouple response style from retrieval.

4.5 We Find No Usable Within-Model Selection Signal in the Features We Tested

Within one model’s candidate set, minority correct values were indistinguishable from minority wrong ones on every feature we measured (votes, dispersion, trace features, log-probabilities; base rate 8%), which is why judges, confidence weighting, and consensus schemes failed in our experiments. What did extract signal is *disagreement between prompts of different style*: on the ~ 25 rows where the recall and grounding prompts disagree by more than 5%, one aggregated decision from an unrelated-family 8B model resolves the tie, giving hasArea $78 \rightarrow 80$ and hasCapacity $26 \rightarrow 28$ on validation; among the four hasArea rows it changed, arbitration corrected three and harmed one. Controls locate the conditions: same-style prompt pairs produce too little disagreement to arbitrate; promoting minority clusters by the arbiter’s mode is *negative*—an 8B cannot overrule a 24B mode, only break ties; and a weak arbiter (Ministral-8B, same family, mode accuracy 23 vs. Llama’s 53 of 100) gains nothing. The isolated v3→v4 ablation—four capacity rows changed, nothing else—confirms the transfer on test: capacity F1 $0.204 \rightarrow 0.214$.

4.6 Which Validation Gains Transfer

Our submissions give per-component test measurements (Table 4). The components that transferred share two properties: each is *gated*, touching only rows that meet an explicit condition, and each gate consults information from outside the model—an unrelated-family vote, or the gold alias structure that makes two surface forms one count. The exemplar effect is the apparent exception: it changes what the backbone retrieves rather than adjudicating between candidates.

The components that did not transfer are fits to validation row distributions: the best of six capacity configurations (winner’s curse), an increment to an already-saturated arbiter, and a city pipeline whose +11-row gain rested on the rate at which split-vote rows happen to be correct-mode. The most instructive case is a tolerance-geometry rescoring we deployed in v9–v10. Two values whose ratio is below $1.05/0.95 \approx 1.105$ admit a single prediction correct for either, so each answer can be shifted inside its own $\pm 5\%$ window to also cover a runner-up cluster. It follows from the metric geometry, has no tuned constants, and gained +1 to +2 rows across every prompt pairing and arbiter seed we tested, yet it cost a row on test (hasArea $0.800 \rightarrow 0.790$), because it moves *every* row rather than only rows meeting a condition; we removed it in v11. Later gated experiments failed the second criterion the same way: certificates computed from the same model’s internal consistency (unit agreement, answer-blind priming) admitted confident errors and lost rows. A 100-row validation set may not reliably expose this form of overfitting; in our experiments, whether a component is gated and whether its gate looks outside the model predicted transfer better than effect size or the presence of a mechanism-level rationale.

4.7 Residual Errors Look Like Values Learned From the Corpus

The residual errors we examined are not reasoning failures. Corfu (gold 626 km²) is answered as 592.9 by every Mistral-family instruct configuration, 593 by Qwen3.5, and 580 by Gemma: stable, model-specific wrong values that survive every prompt. Three observations support the hypothesis that they were learned from training data. First, *turnover across generations*: Qwen3.6 repairs 7 of Qwen3.5’s wrong rows and breaks 11 new ones with nothing else changed, and the union oracle

across generations reaches 89/100. Second, *unrelated training lineages solve mainstream-broken rows*: llm-jp-4-32b and Aya Expanse are far weaker overall yet each partially solves Corfu, which no mainstream-corpus model in our set solves. Third, *the instruct-plus-chat setting is associated with a stronger incorrect anchor*: under raw completion Mistral 3.1 Base scores 10/12-any and 7/12-mode against its own instruct’s 9/4, the gap replicates on Llama-3.1-8B (mode accuracy 60 vs. 53 of 100), and the largest instance is the Gemma lineage of Sec. 4.1. These comparisons change weights and prompting format together; what they establish is that the accessible ceiling is a property of the backbone and its prompting format, not of prompt engineering within one setting.

5 Related Work

Probing factual knowledge in LMs goes back to LAMA (Petroni et al., 2019); the LM-KBC series (Singhanian et al., 2023; Kalo et al., 2024) extends it to KB construction with unconstrained cardinalities. Self-consistency (Wang et al., 2023) aggregates sampled answers by majority vote and ReWiSe (Albert-Roulhac and Zouaq, 2025) tuned it relation-wise; we extend the line to empty answer sets, numeric tolerance, surface-form folding, and cross-family arbitration. Sprague et al. (2025) find chain-of-thought helps reasoning but not knowledge recall; we refine this for recall tasks, where exemplar style moves retrieval although procedures and instructions do not. Ma and Hewitt (2026) report prompting moves parametric knowledge access by 1–2% and RL fine-tuning by $\sim 10\%$; our exemplar effect (+6% on the affected relation) sits between, and fine-tuning is prohibited here. Yuchi et al. (2026) find probe-recoverable numeric knowledge exceeds what models verbalize, one reading of our oracle-mode gaps; Kalyan et al. (2021) and Park et al. (2022) document the magnitude and unit fragilities our decimal-rule and unit-merge components target; and grounding instructions instantiate “According to...” cues (Weller et al., 2024).

6 Conclusion

Our final system is a single pretrained-only 27B backbone queried in a completion format with gated aggregation (validation 0.650, test 0.6493); the two-prompt, arbiter-ensembled 24B system it replaced scores 0.6247. The most reliable finding behind that progression is that backbone recall reli-

ability, together with a prompting format that can reach it, outweighed every prompt-level intervention we tested, and that a half-hour probe ranked candidates well enough to act on.

Two rules earned their place. Choose exemplars, not instructions, to change what a model retrieves; instruction text in our experiments moved only the rejection of already-generated candidates. And before deploying what a small validation set endorses, ask whether the component is gated and whether its gate consults anything outside the model—those that failed this test cost us rows even with a clean mechanism-level rationale. Both rules are bounded by what we measured: three backbones and one task edition.

Limitations

The mechanism findings rest on two model families in depth (Qwen, Mistral) with single-seed generation for most prompts; seed replications exist only for the exemplar, arbitration, and rescoring results. The final Gemma-pt system is the least ablated configuration here: built and submitted within one day, its per-component contributions are not separated on test and its one regression (city, 0.510 $\rightarrow$ 0.490) is unexplained. The 12-row probe is stratified on one system’s failures, and its ranking is validated on the deployment decisions we took, not a held-out model set. Oracle and mode counts on 100-row relations carry ± 1 –2 row noise, which bounds several negative results. Every base-versus-instruct and pretrained-versus-instruct comparison changes weights and prompting format together, which we flag rather than resolve, and the account of residual errors is observational. Aggregation thresholds are validation-tuned, and the transfer criterion of Sec. 4.6 comes from one task edition. Finally, submissions v6–v9 included a second 8B arbiter that put the ensemble at ≈ 39.8 B, over the 32B cap; the violation was ours, found in a self-audit, and v10 removed it at no test cost. Their test numbers should be read as at-risk under the cap rule; v11 and v12 are compliant.

References

Edouard Albert-Roulhac and Amal Zouaq. 2025. [ReWiSe: Relation-wise self-consistency for LLM probing](#). In *Joint Proceedings of KBC-LM and LM-KBC @ ISWC 2025*, volume 4041 of *CEUR Workshop Proceedings*.

Jan-Christoph Kalo, Tuan-Phong Nguyen, Simon

Razniewski, and Bohui Zhang. 2024. [Preface: LM-KBC challenge 2024](#). In *Joint Proceedings of the KBC-LM Workshop and the LM-KBC Challenge @ ISWC 2024*, volume 3853 of *CEUR Workshop Proceedings*, Baltimore, USA.

Ashwin Kalyan, Abhinav Kumar, Arjun Chandrasekaran, Ashish Sabharwal, and Peter Clark. 2021. How much coffee was consumed during EMNLP 2019? Fermi problems: A new reasoning challenge for AI. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.

Melody Ma and John Hewitt. 2026. Improving parametric knowledge access in reasoning language models. [ArXiv:2602.22193](#).

Sungjin Park, Seungwoo Ryu, and Edward Choi. 2022. Do language models understand measurements? In *Findings of the Association for Computational Linguistics: EMNLP 2022*.

Fabio Petroni, Tim Rocktäschel, Sebastian Riedel, Patrick Lewis, Anton Bakhtin, Yuxiang Wu, and Alexander Miller. 2019. Language models as knowledge bases? In *Proceedings of EMNLP-IJCNLP 2019*, pages 2463–2473.

Sneha Singhania, Jan-Christoph Kalo, Simon Razniewski, and Jeff Z. Pan. 2023. [Preface: LM-KBC challenge 2023](#). In *Joint Proceedings of the KBC-LM Workshop and the LM-KBC Challenge @ ISWC 2023*, volume 3577 of *CEUR Workshop Proceedings*, Athens, Greece.

Zayne Sprague, Fangcong Yin, Juan Diego Rodriguez, Dongwei Jiang, Manya Wadhwa, Prasann Singhal, Xinyu Zhao, Xi Ye, Kyle Mahowald, and Greg Durrett. 2025. To CoT or not to CoT? chain-of-thought helps mainly on math and symbolic reasoning. In *The Thirteenth International Conference on Learning Representations (ICLR)*.

Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-consistency improves chain of thought reasoning in language models. In *The Eleventh International Conference on Learning Representations (ICLR)*.

Orion Weller, Marc Marone, Nathaniel Weir, Dawn Lawrie, Daniel Khashabi, and Benjamin Van Durme. 2024. “according to ...”: Prompting language models improves quoting from pre-training data. In *Proceedings of EACL 2024*.

Fengting Yuchi, Li Du, and Jason Eisner. 2026. LLMs know more about numbers than they can say. In *Proceedings of EACL 2026 (Short Papers)*. [ArXiv:2602.07812](#).